

# The space time complexity

## Problem statement

It is year 2031 and a collective consisting of leading scientists on the earth called TSF (*The space frontiers*) has open sourced the design for a nuclear powered space probe. It is tiny in size, can be 3D printed and assembled, launched with specialized [propellants](#) from anywhere and has a really long life span. It makes it an ideal candidate to send it for non critical data collection missions into deep-space. This has given rise to a movement **#mapOurUniverse**. TSF is crowdsourcing the effort to launch millions of probes to help map the entire universe. These probes pair up to create a mesh and each probe is configured to send data packets of their environment and location to a server back on the earth.

You have been set incharge to create a **critical**, performant and resilient **data store from scratch** that can store and maintain the latest entry sent by each probe at any given time. This is a crucial low level component for building a map of the universe as it records the farthest known edge of the universe at any given time. Several other systems query & act on data collected in this system in real time to adjust probe flight path.

Following are the demands from the system to work as expected

- The system should expose an HTTP endpoint to accept a PUT request containing a payload specified below in the document
- The system should create/update the payload based on a unique probe id. The spec for probe id is specified below in the document
- The system should not lose data, if it undergoes a planned (an explicit process SIGTERM/SIGQUIT) or an unplanned (*due to a failure, SIGKILL*) restart
- The system needs to have only the latest record per probe. Older records can be overwritten safely.
- The system needs to expose an HTTP endpoint to retrieve the latest record by a probe id
- The 99p response time for writing a record to the service should not go beyond 200ms.
- The probes may have unreliable gaps between two subsequent transmissions. This should not impact the availability of their last available record
- The system should also allow data to be queried from the system by a probe id. The contract and nature of data queries are specified in the section below
- The system should identify latest event using eventTransmissionTime and discard older events in case they are received in an incorrect sequence (*an older event after the latest one*)

## Probe Id spec

An alphanumeric id. It has to be unique for a probe. A probe can have only one probe id. At Least 3 characters and maximum 100.

## Payload spec

Average message size - 5 kb (Min 1 kb to max 20 kb)

Field name	Description	Value format
probeId	A unique id per probe hardwired into the probe sent in each transmission	Spec for probe id specified below
eventId	A unique id per event.	UUID with timestamp built in with millisecond precision
messageType	A message classification. For this scenario it will always be "spaceCartography"	Fixed string
eventReceivedTime	Timestamp when the event information was received. Needs to be populated by the receiving applications/datastores	Milliseconds since epoch
eventTransmissionTime	Timestamp when the event was transmitted from the probe. Injected into the payload by the probe	Milliseconds since epoch
messageData.measureName	Name of a measurement. Total 6 types of measurements exist	Fixed string set
messageData.measureCode	A short code for a measurement. Total 6 types of measurement code exist  <b>SCSED</b> - The euclidean distance from earth to	Fixed String set

	<p>the probe as per the specifications of Spherical coordinate system</p> <p><b>SCSEAA</b> - The azimuth angle from earth to the probe as per the specifications of <a href="#">Spherical coordinate system</a></p> <p><b>SCSEPA</b> - The polar angle from earth to the probe as per the specifications of Spherical coordinate system. Measured in <a href="#">parsecs</a></p> <p><b>LER</b> - Frequency of electromagnetic radiation in the current environment / space</p> <p><b>PLSE</b> - Remaining life span for the probe in number of years accounting current wear and tear</p> <p><b>PDL</b> - Diagnostic logs from the probe</p>	
messageData.measureUnit	Unit of measurement for the specified measurement	Fixed String set
messageData.measureValue	The actual value derived/raw for the measurement	<p><b>SCSED</b> - Floating point</p> <p><b>SCSEAA</b> - Floating point</p> <p><b>SCSEPA</b> - Floating point</p> <p><b>LER</b> - Floating point</p> <p><b>PLSE</b> - Floating point</p> <p><b>PDL</b> - Text</p>

messageData.measureValueDescription	Description of the measurement	Fixed string set
messageData.measureType	<p>Category classification for the measurement. Three categories exist currently.</p> <p><b>Positioning</b> - The position of the probe in space in relation to the position of Earth.</p> <p><b>Composition</b> - The composition of the space around the probe</p> <p><b>Probe</b> - Details about the probe itself</p>	Fixed string set
messageData.componentReading	The raw reading from the underlying component in case the measured value is derived	Floating point

### Sample Payload

```
{
  "probeId": "PRB34222422421123",
  "eventId": "7707d6a0-61b5-11ec-9f10-0800200c9a66",
  "messageType": "spaceCartography",
  "eventReceivedTime": <TO BE POPULATED BY DATA STORE AS EPOCH WITH MILLIS>,
  "eventTransmissionTime": 1640018265951,
  "messageData": [
    {
      "measureName": "Spherical coordinate system - euclidean distance",
      "measureCode": "SCSED",
      "measureUnit": "parsecs",
      "measureValue": 5399e5,
      "measureValueDescription": "Euclidean distance from earth",
      "measureType": "Positioning",
      "componentReading": 43e23
    },
    {
```

```

        "measureName": "Spherical coordinate system - azimuth
angle",
        "measureCode": "SCSEAA",
        "measureUnit": "degrees",
        "measureValue": 170.42,
        "measureValueDescription": "Azimuth angle from earth",
        "measureType": "Positioning",
        "componentReading": 46e2
    },
    {
        "measureName": "Spherical coordinate system - polar
angle",
        "measureCode": "SCSEPA",
        "measureUnit": "degrees",
        "measureValue": 30.23,
        "measureValueDescription": "Polar/Inclination angle from
earth",
        "measureType": "Positioning",
        "componentReading": 56e42
    },
    {
        "measureName": "Localized electromagnetic frequency
reading",
        "measureCode": "LER",
        "measureUnit": "hz",
        "measureValue": 3e5,
        "measureValueDescription": "Electromagnetic frequency
reading",
        "measureType": "Composition",
        "componentReading": 3e15
    },
    {
        "measureName": "Probe lifespan estimate",
        "measureCode": "PLSE",
        "measureUnit": "Years",
        "measureValue": 2390e2,
        "measureValueDescription": "Number of years left in
probe lifespan",
        "measureType": "Probe",
        "componentReading": 6524e3
    },
    {
        "measureName": "Probe diagnostic logs",
        "measureCode": "PDL",
        "measureUnit": "Text",
        "measureValue": "some log data from probe",
        "measureValueDescription": "the diagnostic information
from the probe",

```

```

        "measureType": "Probe",
        "componentReading": 0.0
    }
]
}

```

The payload can be stored in any format as long as it can be constructed back and returned as provided.

## Probe data payload contract

REST **PUT** /probe/<probe\_id>/event/<event\_id>

*This endpoint will be used by the probe simulator to submit packets and test the solution throughput. If the probe id doesn't exist then a new record should be created for the probe with respective event/event id.*

## Probe query contract

REST **GET** /probe/<probe\_id>/latest

*This endpoint will be used by the probe simulator to test the availability of the data that was posted previously and the time within which it is available*

## Body

As specified in payload spec above

## Evaluation criteria

Criteria	Point allocation formula
99p response time for saving a record in data store < 200 ms	
Each probe will transmit a message containing above payload every 10 seconds	
Points based on peak load factor sustained for 120 minutes  Load factor = <i>Numbers of probes supported within thresholds specified above.</i>	X  *where X = sustained peak probe count
Records written to the data store > 5000 ms should not be lost during a planned/unplanned restart. For durability	$X * \left( 4 \left( 1 - \left( \frac{\text{guaranteed-durability-millis}}{5000} \right) \right) \right)$

<p>guarantees less than 5000 ms, there will be bonus points as per the formula specified here. For durability guarantee more than 5000 ms there will be a penalty for the points as demonstrated in the formula here</p>	<p><i>*where X = sustained peak probe count with claimed durability duration</i></p> <p>For instance</p> <p>A peak probe load of 3000 probes with a guaranteed durability after 100 ms will be awarded 11671 points.</p> <p>A peak probe load of 10000 probes with a guaranteed durability after 4900 ms will be awarded 10281 points.</p> <p>A peak probe load of 30000 probes with a guaranteed durability after 10000 ms will be awarded 7500 points.</p> <p>A peak probe load of 100000 probes with no guaranteed durability will be awarded 0 points.</p>
--	--

## Guidelines and constraints

- Once the problem statement is released you can start developing a solution. Specific hackathons may be arranged to carve out time to enable people to build solutions, but the time is not restricted to that.
- Any runtime/technology can be used to develop the data store itself. As long as the fundamental storage capability is not pre-built. E.g. using an existing storage engine.
- You cannot use a pre-existing data store (Relation, NOSQL or any other type). The data store should be hand crafted, custom-built and tuned.
- Given the test scripts will expect an HTTP api, either the database can expose an HTTP api directly or it can be fronted by a web service to channel Reads/Writes
- No intermediary caches or storage side cars are allowed, in case the data store is fronted by a web service. This is different from in memory indexes that you may create for the data store.
- Everyone will be provided with the exact same hardware. The machines are likely to be Ubuntu **c6gd.large** (4 GiB, 1 x 118 NVMe SSD, Up to 10 Gigabit) machines in AWS Mumbai region. This instance comes with a NVMe SSD attached to the instance. The default capacity of the SSD will be flat 20GB for all submissions. However once an implementation crosses 1 million concurrent probes, for each subsequent million probes an additional 20 GB will be allocated.
- There will be an external script that generates payload with the spec specified in the document. The devised solution should be able to interoperate with the external script
- The cloud machines will be provisioned after the requested intermediary stats have been published from local development. Details on intermediary stats will be available once the competition kicks off.

## Evaluation process

- Big O kicks off (Overall duration 8 weeks)
- Each team starts building data store locally on their machines
- Intermediary stats are requested. Teams that publish the stats get the next level of resources at each step.
  - Access to AWS
  - Access to test scripts
  - Access to final testing setup and so on.
- Sessions and hackathons are organized to provide ideas, knowledge, tips/tricks & pizzas
- Each team is able to test on AWS with final scripts and measure their storage performance
- Evaluation setup is created and executed for all qualifying teams
- The top two teams with highest sustained load factor are declared as winners

## Hackathon schedule

Date	Event / Location
<del>20th Dec 2021</del>	<del>Kick off event with email requesting registration</del>
<del>22nd Dec 2021</del>	<del>Broadcast problem statement to the registered teams</del>
<del>24th Dec 2021</del>	<del>Close registrations</del>
<del>7th Jan 2021</del>	<del>First hackathon</del>
21st Jan 2021	Second hackathon
4th Feb 2021	Third hackathon
18th Feb 2021	Midnight submission close
18th Feb 2021 - 28th Feb 2021	Daily showdowns and showcases (one team at a time)
Annual day	Winner declaration and prize distribution